

Unit5HW

May 19, 2020

1 Unit5作业

```
[18]: from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = 'all'
import scipy.stats as stats
import numpy as np
import statsmodels.stats.proportion as proportion
import statsmodels.formula.api as smf
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
sns.set_style('darkgrid')
```

HW-U5-1: 对于Diabetes.csv数据, 请利用协方差和pearson 相关系数分析Glucose与Blood-Pressure的关系 (0.5分), 并画出血糖-血压散点图 (0.5分)

```
[19]: diabetes=pd.read_csv('Diabetes.csv')
diabetes.head(5)
```

```
[19]:
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[20]: diabetes=diabetes[diabetes['Glucose']!=0]
diabetes=diabetes[diabetes['BloodPressure']!=0]
glu=diabetes['Glucose']
bp=diabetes['BloodPressure']
```

```
plt.scatter(glu,bp)
```

[20]: <matplotlib.collections.PathCollection at 0x1f489300508>



```
[21]: np.cov(glu,bp)
stats.pearsonr(glu,bp)
```

[21]: array([[941.21371888, 84.81198513],
[84.81198513, 153.4157062]])

[21]: (0.223191778249542, 1.138581203805524e-09)

pearson' $s_r = 0.22, p = 1.14 \times 10^{-9}$. 该结果显示Glucose与BloodPressure之间没有明显的相关关系。

HW-5-2: 对于Titanic.csv数据（参见前面单元作业），请用分别用pearson r, spearman rho, kendall' s tau分别计算乘客年龄与买的票的等级的相关系数（1分）

```
[22]: titantic=pd.read_csv('Titanic.csv')
titantic.head(5)
titantic=titantic.dropna()
```

[22]:

	Name	PClass	Age	Sex	\
0	Allen, Miss Elisabeth Walton	1st	29.00	female	
1	Allison, Miss Helen Loraine	1st	2.00	female	
2	Allison, Mr Hudson Joshua Creighton	1st	30.00	male	
3	Allison, Mrs Hudson JC (Bessie Waldo Daniels)	1st	25.00	female	

```
4 Allison, Master Hudson Trevor 1st 0.92 male
```

```
Survived
0      1
1      0
2      0
3      0
4      1
```

```
[23]: age=titanic['Age']
      pclass=titanic['PClass']
      pclass=pclass.replace(['1st', '2nd', '3rd'], [1,2,3])
```

```
[24]: stats.pearsonr(age,pclass)
      stats.spearmanr(age,pclass)
      stats.kendalltau(age,pclass)
```

```
[24]: (-0.4141214595264922, 1.0969903610990536e-32)
```

```
[24]: SpearmanrResult(correlation=-0.39366216507025165, pvalue=1.9756955531661058e-29)
```

```
[24]: KendalltauResult(correlation=-0.3103724477828569, pvalue=8.600565143832718e-28)
```

以上三个结果的p-value均很小，均反映了Age与PClass之间没有显著的相关关系。

HW-5-3：针对汽车数据mtcars.csv

(1) 画出wt~ mpg散点图；用简单线性回归分析mpg（因变量），和wt（自变量）的关系，并根据回归结果中的截距和斜率及其显著性水平（p值），对结果进行解释；并解释Rsquare. (1.5分)


```
[25]: df=pd.read_csv('mtcars.csv')
      df.head(5)
```

```
[25]: Unnamed: 0  mpg  cyl  disp  hp  drat    wt   qsec  vs  am  gear  \
0      Mazda RX4  21.0    6  160.0  110  3.90  2.620  16.46  0   1     4
1      Mazda RX4 Wag  21.0    6  160.0  110  3.90  2.875  17.02  0   1     4
2      Datsun 710   22.8    4  108.0   93  3.85  2.320  18.61  1   1     4
3      Hornet 4 Drive  21.4    6  258.0  110  3.08  3.215  19.44  1   0     3
4      Hornet Sportabout  18.7    8  360.0  175  3.15  3.440  17.02  0   0     3
```

```
carb
0      4
1      4
2      1
3      1
4      2
```

```
[26]: sns.scatterplot(x='wt',y='mpg',data=df)
```

```
[26]: <matplotlib.axes._subplots.AxesSubplot at 0x1f489864948>
```



```
[27]: result1=smf.ols('mpg~wt',data=df).fit()
result1.summary()
```

```
[27]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                        OLS Regression Results
=====
Dep. Variable:          mpg      R-squared:                0.753
Model:                  OLS      Adj. R-squared:           0.745
Method:                 Least Squares      F-statistic:        91.38
Date:                   Tue, 19 May 2020     Prob (F-statistic):    1.29e-10
Time:                   20:04:26      Log-Likelihood:       -80.015
No. Observations:       32      AIC:                  164.0
Df Residuals:           30      BIC:                  167.0
Df Model:                1
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	37.2851	1.878	19.858	0.000	33.450	41.120
wt	-5.3445	0.559	-9.559	0.000	-6.486	-4.203

```

=====
```

```
=====
Omnibus:                2.988    Durbin-Watson:                1.252
Prob(Omnibus):          0.225    Jarque-Bera (JB):          2.399
Skew:                   0.668    Prob(JB):                  0.301
Kurtosis:               2.877    Cond. No.                  12.7
=====
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
"""
```

回归结果 $mpg = -5.3445wt + 37.2851$, 截距和斜率的p值均很小, 说明截距和斜率均有显著意义. $R^2 = 0.753$ 说明mpg与wt之间有比较强的线性关系(strong).

(2) 用多元线性回归分析mpg (因变量), 和wt(自变量)、hp(自变量)的关系, 并根据回归结果中的各个自变量的系数及其显著性水平 (p值), 对回归结果进行解释; 并解释R-square (1.5分)

```
[28]: result2=sfm.ols('mpg~wt+hp',data=df).fit()
result2.summary()
```

```
[28]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                OLS Regression Results
=====
Dep. Variable:                mpg    R-squared:                0.827
Model:                        OLS    Adj. R-squared:          0.815
Method:                        Least Squares    F-statistic:            69.21
Date:                        Tue, 19 May 2020    Prob (F-statistic):      9.11e-12
Time:                        20:04:26    Log-Likelihood:          -74.326
No. Observations:              32    AIC:                    154.7
Df Residuals:                  29    BIC:                    159.0
Df Model:                      2
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	37.2273	1.599	23.285	0.000	33.957	40.497
wt	-3.8778	0.633	-6.129	0.000	-5.172	-2.584
hp	-0.0318	0.009	-3.519	0.001	-0.050	-0.013

```
=====
Omnibus:                5.303    Durbin-Watson:                1.362
Prob(Omnibus):          0.071    Jarque-Bera (JB):          4.046
Skew:                   0.855    Prob(JB):                  0.132
Kurtosis:               3.332    Cond. No.                  588.
=====
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
"""
```

回归结果 $mpg = -3.8778wt - 0.0318hp + 37.2273$ ，三个系数的p值均小于0.05，说明结果具有显著意义。mpg与wt和hp都成负相关，但且wt对mpg的影响要高于hp。

$R^2 = 0.827$ ，说明mpg与wt和hp之间有很强的二元线性关系。相比单变量回归， R^2 值提高了，结合 $Adj-R^2$ ，也没有出现明显的过拟合现象，因此模型较单变量效果更好。

HW-5-4: mtcars.csv 数据)，汽车的离合 (am: 手动 (1) /自动 (0)) 与汽车的油耗 (mpg), 马力 (hp)是很相关的，请：

(1) 基于全部数据用mpg,hp作为自变量，am作为因变量，建立对应的逻辑回归模型，并作出解释 (1.0分)

```
[29]: log_res=smf.logit('am~mpg+hp',data=df).fit()
log_res.summary()
```

```
Optimization terminated successfully.
      Current function value: 0.300509
      Iterations 9
```

```
[29]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                        Logit Regression Results
=====
Dep. Variable:          am      No. Observations:          32
Model:                Logit      Df Residuals:             29
Method:                MLE       Df Model:                2
Date:                  Tue, 19 May 2020      Pseudo R-squ.:       0.5551
Time:                  20:04:26      Log-Likelihood:       -9.6163
converged:              True      LL-Null:             -21.615
Covariance Type:        nonrobust      LLR p-value:         6.153e-06
=====
               coef      std err          z      P>|z|      [0.025      0.975]
-----
Intercept    -33.6052     15.077     -2.229     0.026    -63.156    -4.055
mpg           1.2596      0.567      2.220     0.026      0.147      2.372
hp            0.0550      0.027      2.045     0.041      0.002      0.108
=====
```

```
Possibly complete quasi-separation: A fraction 0.12 of observations can be
perfectly predicted. This might indicate that there is complete
quasi-separation. In this case some parameters will not be identified.
"""
```

mpg和hp的p值均小于0.05，说明这两者与am均有关联，且mpg在其中占的比重更大。

附加题：（2）将前20条记录作为训练数据，重新建立上面的逻辑回归模型，然后用后12条记录作为测试数据，再对该模型进行测试，并对结果作出解释。

```
[34]: from sklearn.model_selection import train_test_split
train, test = train_test_split(df, test_size=12/32)
###Normalization
# Assuming same lines from your example
cols_to_norm = "mpg", "hp"
train[cols_to_norm] = train[cols_to_norm].apply(lambda x: (x - x.mean()) / (x.
↪std()))
test[cols_to_norm] = test[cols_to_norm].apply(lambda x: (x - x.mean()) / (x.
↪std()))
```

```
[31]: model=smf.logit('am~mpg+hp',data=train).fit()
model.summary()
```

Optimization terminated successfully.

Current function value: 0.210510

Iterations 10

```
[31]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                        Logit Regression Results
=====
Dep. Variable:          am    No. Observations:          20
Model:                Logit    Df Residuals:           17
Method:                MLE     Df Model:              2
Date:                 Tue, 19 May 2020    Pseudo R-squ.:       0.6941
Time:                 20:04:26    Log-Likelihood:      -4.2102
converged:              True    LL-Null:           -13.763
Covariance Type:      nonrobust    LLR p-value:        7.102e-05
=====
               coef      std err          z      P>|z|      [0.025      0.975]
-----
Intercept         1.8711      1.363      1.373      0.170      -0.800      4.542
mpg              11.8868      6.970      1.705      0.088      -1.774     25.547
hp                 5.6484      3.435      1.644      0.100      -1.084     12.381
=====
```

Possibly complete quasi-separation: A fraction 0.30 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

```
"""
```

```
[32]: trainingRes=pd.DataFrame(model.pred_table())
trainingRes.columns=["Predicted Outcome 0","Predicted Outcome 1"]
```

```

trainingRes=trainingRes.rename(index={0:"Actual Outcome 0", 1:"Actually Outcome_
↪1"})
trainingRes
print("The accuracy for the train data is :", (trainingRes.iloc[1,1]+trainingRes.
↪iloc[0,0])/20)

```

```

[32]:
          Predicted Outcome 0  Predicted Outcome 1
Actual Outcome 0              8.0              1.0
Actually Outcome 1            1.0             10.0

```

The accuracy for the train data is : 0.9

```

[33]: pred_values = model.predict(test)
bins=np.array([0,0.5,1])
cm = np.histogram2d(test['am'], pred_values, bins=bins)[0]
accuracy = (cm[0,0]+cm[1,1])/cm.sum()
print("The prediction accuracy for the test data is :", accuracy)

testRes=pd.DataFrame(cm)
testRes.columns=["Predicted Outcome 0","Predicted Outcome 1"]
testRes=testRes.rename(index={0:"Actual Outcome 0", 1:"Actually Outcome 1"})
testRes

```

The prediction accuracy for the test data is : 0.6666666666666666

```

[33]:
          Predicted Outcome 0  Predicted Outcome 1
Actual Outcome 0              6.0              4.0
Actually Outcome 1            0.0              2.0

```

从拟合的模型来看，由于训练集较之前规模减小($32 - > 20$)，因此模型精确度有所下降，三个p值均较大，可能存在一定的欠拟合情况。但从结果来看模型效果还不错，训练集的正确率为0.8，测试集的准确率有0.917。